

MASSACHUSETTS INSTITUTE OF TECHNOLOGY
DEPARTMENT OF ELECTRICAL ENGINEERING AND COMPUTER SCIENCE

PROPOSAL FOR THESIS RESEARCH IN PARTIAL FULFILLMENT
OF THE REQUIREMENTS FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

TITLE: Protecting User Data in Large-Scale Web Services

SUBMITTED BY: Frank Wang
32 Vassar Street, #32-G978
Cambridge, MA 02139

(SIGNATURE OF AUTHOR)

DATE OF SUBMISSION: May 2, 2018
EXPECTED DATE OF COMPLETION: June 2018
LABORATORY: Computer Science and Artificial Intelligence Laboratory

BRIEF STATEMENT OF THE PROBLEM:

Web services like Google, Facebook, and Dropbox are now an essential part of peoples lives. Users willingly provide their data to these services because these services deliver substantial value in return through their centralization and analysis of data, such product recommendations and ability to easily share information. To provide this value, these services collect, store, and analyze large amounts of their users sensitive data. However, once the user provides her information to the web service, she loses control over how the application manipulates that data. For example, a user cannot control where the application forwards her data. Even if the service wanted to allow users to define access controls, it is unclear how these access controls should be expressed and enforced. Not only is it difficult to develop these secure access control mechanisms, but it is also difficult to ensure these mechanisms are practical. My thesis addresses these concerns.

1 Introduction

Web services like Google, Facebook, and Dropbox are now an essential part of peoples lives. Users willingly provide their data to these services because these services deliver substantial value in return through their centralization and analysis of data, such product recommendations and ability to easily share information. For example, users are willing to share their data with Facebook to learn about the social lives of their friends as well as share their own social lives more easily. Similarly, users are willing to provide their data to Amazon to discover better product recommendations. To provide value to users, these services collect, store, and analyze large amounts of their users' sensitive data. However, once the user provides her information to the web service, she *loses control* over how the application manipulates that data. For example, a user cannot control where the application forwards her data. Even if the service wanted to allow users to define access controls, it is unclear how these access controls should be expressed and enforced. Not only is it difficult to develop these secure access control mechanisms, but it is also difficult to ensure these mechanisms are *practical*. This thesis addresses these concerns. More specifically, it focuses on *building practical, secure mechanisms for protecting user data in large-scale, distributed web services*.

In this thesis, I will describe three systems that address a variety of concerns around data leakage in web applications. The first two systems focus on protecting user data against server-side leakage. Splinter leverages a recent cryptographic primitive, function secret sharing, to practically execute these queries without revealing sensitive information to the servers. The next system, Riverbed, provides practical information flow control for distributed systems without requiring developers to label state or write code in special languages. The final system, Veil, focuses on client-side leakage. Veil allows web page developers to enforce stronger private browsing semantics without browser support.

2 Splinter

2.1 Motivation

Many online services let users query large public datasets: some examples include restaurant sites, product catalogs, stock quotes, and searching for directions on maps. In these services, any user can query the data, and the datasets themselves are not sensitive. However, web services can infer a great deal of identifiable and sensitive user information from these queries, such as her current location, political affiliation, sexual orientation, income, etc. [1, 2]. Web services can use this information maliciously and put users at risk to practices such as discriminatory pricing [3, 4, 5]. For example, online stores have charged users different prices based on location [6], and travel sites have also increased prices for certain frequently searched flights [7]. Even when the services are honest, server compromise and subpoenas can leak the sensitive user information on these services [8, 9, 10].

This thesis presents Splinter, a system that protects users' queries on public datasets while achieving practical performance for many current web applications. In Splinter, the user divides each query into shares and sends them to different *providers*, which are services hosting a copy of the dataset (Figure 1). As long as any one of the providers is honest and does not collude with the others, the providers cannot discover sensitive information in the query. However, given responses from all the providers, the user can compute the answer to her query.

Previous private query systems have generally not achieved practical performance because they use expensive cryptographic primitives and protocols. For example, systems based on Private Information Retrieval (PIR) [11, 12, 13] require many round trips and high bandwidth for complex queries, while systems based on garbled circuits [14, 15, 16] have a high computational cost. These approaches are especially costly for mobile clients on high-latency networks.

Instead, Splinter uses and extends a recent cryptographic primitive called Function Secret

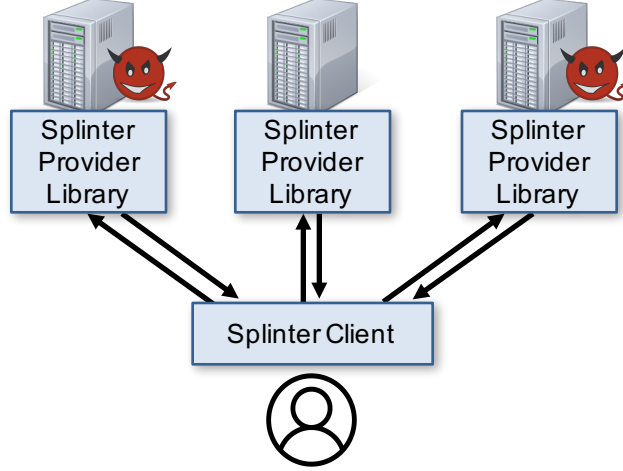


Figure 1: Splinter architecture. The Splinter client splits each user query into shares and sends them to multiple providers. It then combines their results to obtain the final answer. The user’s query remains private as long as any one provider is honest.

Sharing (FSS) [17, 18], which makes it up to an order of magnitude faster than prior systems. FSS allows the client to split certain functions into shares that keep parameters of the function hidden unless all the providers collude. With judicious use of FSS, many queries can be answered at low CPU and bandwidth cost in only a single network round trip.

Splinter makes two contributions over previous work on FSS. First, prior work has only demonstrated efficient FSS protocols for point and interval functions with additive aggregates such as SUMs [17]. We present protocols that support a more complex set of non-additive aggregates such as MAX/MIN and TOPK at low computational and communication cost. Together, these protocols let Splinter support a subset of SQL that can capture many popular online applications.

Second, we develop an optimized implementation of FSS for modern hardware that leverages AES-NI [19] instructions and multicore CPUs. For example, using the one-way compression functions that utilize modern AES instruction sets, our implementation is $2.5\times$ faster per core than a naïve implementation of FSS. Together, these optimizations let Splinter query datasets with millions of records at sub-second latency on a single server.

We evaluate Splinter by implementing three applications over it: a restaurant review site similar to Yelp, airline ticket search, and map routing. For all of our applications, Splinter

can execute queries in less than 1.6 seconds, at a cost of less than 0.02¢ in server resources on Amazon EC2. Splinter’s low cost means that providers could profitably run a Splinter-based service similar to OpenStreetMap routing [20], an open-source maps service, while only charging users a few dollars per month.

Splinter aims to protect sensitive information in users’ queries from providers. This section provides an overview of Splinter’s architecture, security goals, and threat model.

2.2 Splinter Overview

There are two main principals in Splinter: the *user* and the *providers*. Each provider hosts a copy of the data. Providers can retrieve this data from a public repository or mirror site. For example, OpenStreetMap [20] publishes publicly available map, point-of-interest, and traffic data. For a given user query, all the providers have to run it on the same view of the data. Maintaining data consistency from mirror sites is beyond the scope of this thesis, but standard techniques can be used [21, 22].

As shown in Figure 1, to issue a query in Splinter, a user splits her query into *shares*, using the Splinter client, and submits each share to a different provider. The user can select any providers of her choice that host the dataset. The providers use their shares to execute the user’s query over the cleartext public data, using the Splinter provider library. As long as one provider is *honest* (does not collude with others), the user’s sensitive information in the original query remains private. When the user receives the responses from the providers, she combines them to obtain the final answer to her original query.

2.3 Security Goals

The goal of Splinter is to hide sensitive parameters in a user’s query. Specifically, Splinter lets users run *parametrized queries*, where both the parameters and query results are hidden from providers. For example, consider the following query, which finds the 10 cheapest flights between a source and destination:

```
SELECT TOP 10 flightid FROM flights
```

```
WHERE source = ? AND dest = ?  
  
ORDER BY price
```

Splinter hides the information represented by the question marks, i.e., the source and destination in this example. The column names being selected and filtered are not hidden. Finally, Splinter also hides the query’s results—otherwise, these might be used to infer the source and destination. Splinter supports a subset of the SQL language.

The easiest way to achieve this property would be for users to download the whole database and run the queries locally. However, this requires substantial bandwidth and computation for the user. Moreover, many datasets change constantly, e.g., to include traffic information or new product reviews. It would be impractical for the user to continuously download these updates. Therefore, our performance objective is to minimize computation and communication costs. For a database of n records, Splinter only requires $O(n \log n)$ computation at the providers and $O(\log n)$ communication.

2.4 Threat Model

Splinter keeps the parameters in the user’s query hidden as long as at least one of the user-chosen providers does not collude with others. Splinter also assumes these providers are *honest but curious*: a provider can observe the interactions between itself and the client, but Splinter does not protect against providers returning incorrect results or maliciously modifying the dataset.

We assume that the user communicates with each provider through a secure channel (e.g., using SSL), and that the user’s Splinter client is uncompromised. Our cryptographic assumptions are standard. We only assume the existence of one-way functions in our two-provider implementation. In our implementation for multiple providers, the security of Paillier encryption [23] is also assumed.

3 Riverbed

3.1 Motivation

In a web service, a client like a desktop browser or a smartphone app interacts with datacenter machines. Although smartphones and web browsers provide rich platforms for computation, the core application state typically resides in cloud storage. This state accrues much of its value from server-side computations that involve no participation (or explicit consent) from end-user devices.

By running the bulk of an application within VMs in a commodity cloud, developers receive two benefits. First, developers shift the burden of server administration to professional datacenter operators. Second, developers gain access to scale-out resources that vastly exceed those that are available to a single user device. Scale-out storage allows developers to co-locate large amounts of data from multiple users; scale-out computation allows developers to process the co-located data for the benefit of users (e.g., by providing tailored search results) and the benefit of the application (e.g., by providing targeted advertising).

A Loss of User Control: Unfortunately, there is a disadvantage to migrating application code and user data from a user’s local machine to a remote datacenter server: the user loses control over where her data is stored, how it is computed upon, and how the data (and its derivatives) are shared with other services. Users are increasingly aware of the risks associated with unauthorized data leakage [24, 25, 26], and some governments have begun to mandate that online services provide users with more control over how their data is used. For example, in 2016, the EU passed the General Data Protection Regulation [27]. Articles 6, 7, and 8 of the GDPR state that users must give consent for their data to be accessed. Article 15 establishes a “right to access,” mandating that users have the ability to know how their data is being processed by a company. Article 17 defines a user’s right to request her data to be deleted; Article 32 requires a company to implement “appropriate” security measures for data-handling pipelines. Unfortunately, requirements like these *lack strong*

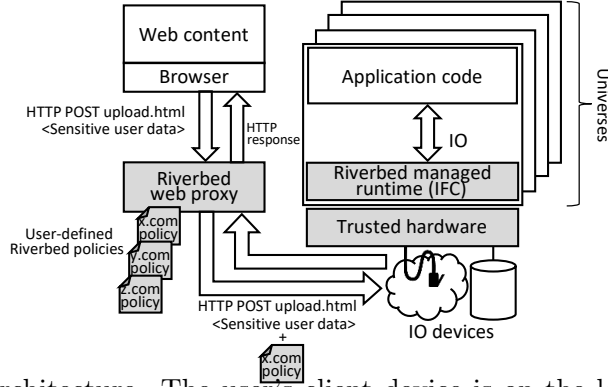


Figure 2: Riverbed’s architecture. The user’s client device is on the left, and the web service is on the right. Unmodified components are white; modified or new components are grey.

definitions and enforcement mechanisms at the systems level. Laws like GDPR provide little technical guidance to a developer who wants to comply with the laws while still providing the sophisticated applications that users enjoy.

The research community has proposed information flow control (IFC) as a way to constrain how sensitive data spreads throughout a complex system [28, 29]. IFC assigns labels to program variables or OS-level resources like processes and pipes; given a partial ordering which defines the desired security relationships between labels, an IFC system can enforce rich properties involving data secrecy and integrity. Unfortunately, traditional IFC is too burdensome to use in modern, large-scale web services. Creating and maintaining a partial ordering of labels is too difficult—the average programmer or end-user struggles to reason about data safety using the abstraction of fine-grained label hierarchies. As a result, no popular, large-scale web service uses IFC to restrict how sensitive data is processed and shared.

3.2 Riverbed Overview

In this thesis, we introduce Riverbed, a distributed web platform for safeguarding the privacy of user data. Riverbed provides benefits to both web developers and to end users. To web developers, Riverbed provides a *practical* IFC system which allows developers to easily “bolt on” stronger security policies for complex applications written in standard managed languages. To end users, Riverbed provides a straightforward mechanism to verify that

server-side code is running within a privacy-preserving environment.

Figure 3 describes Riverbed’s architecture. For each Riverbed web service, a user defines an information flow policy using simple, human-understandable constraints like “do not save my data to persistent storage” or “my data may only be sent over the network to `foo.com`.” In the common case, users employ predefined, templated policy files that are designed by user advocacy groups like the EFF. When the user generates an HTTP request, a web proxy on the user’s device transparently adds the appropriate data flow policy as a special HTTP header.

Within a datacenter, Riverbed leverages the fact that many services run atop managed runtimes like Python, .NET, or the JVM. Riverbed modifies such a runtime to automatically taint incoming HTTP data with the associated user policies. As the application derives new data from tainted bytes, the runtime ensures that the new data is also marked as tainted. The runtime ensures that, whenever an application tries to externalize data via the disk or the network, the externalization is only allowed if it is permitted by user policies. An application process which attempts a disallowed externalization is automatically terminated.

In Riverbed, application code (i.e., the code which the managed runtime executes) is totally unaware that IFC is occurring. Indeed, application developers have no way to read, write, create, or destroy taints and data flow policies. The advantage of this scheme is that it makes Riverbed compatible with code that has not been explicitly annotated with traditional IFC labels. However, different end users will likely define incompatible data flow policies. As a result, policy-agnostic code would quickly generate a policy violation for some subset of users; Riverbed would then terminate the application. To avoid this problem, Riverbed spawns multiple, lightweight copies of the service, one for each set of users who share the same data flow policies. We call each copy a *universe*. Since users in the same universe allow the same types of data manipulations, any policy violations indicate true problems with the application (e.g., the application tried to send sensitive data to a server that was not whitelisted by the inhabitants of the universe). Developers can then fix these violations

to ensure that code respects the desired privacy protections for the target user demographic.

Before a user’s Riverbed proxy sends data to a server, the proxy employs remote software attestation [30] to verify that the server is running an IFC-enforcing Riverbed runtime. Importantly, a trusted server will perform *next-hop attestation*—the server will not transmit user data to another network endpoint unless that endpoint is an attested Riverbed runtime whose TLS certificate name is explicitly whitelisted by the user’s data flow policy. In this manner, Riverbed enables controlled data sharing between machines that belong to different services.

3.3 Our Contributions

To the best of our knowledge, Riverbed is the first distributed IFC framework which is practical enough to support large-scale, feature-rich web services that are written in general-purpose managed languages. Riverbed preserves the traditional advantages of cloud-based applications, allowing developers to offload administrative tasks and leverage scale-out resources. However, Riverbed’s universe mechanism, coupled with a simple policy language, provides users with understandable, enforceable abstractions for controlling how datacenters manipulate sensitive data. Riverbed makes it easier for developers to comply with laws like GDPR—users give explicit consent for data access via Riverbed policies, with server-side universes constraining how user data may be processed, and where its derivatives can be stored.

We have ported several non-trivial applications to Riverbed, and written data flow policies for those applications. Our experiments shows that Riverbed enforces data flow policies with worst-case end-to-end overheads of 10%; Riverbed also supports legacy code with little or no developer intervention, making it easy for well-intentioned (but average-skill) developers to write services that respect user privacy.

3.4 Threat Model

Riverbed assumes that developers want to enforce user-defined privacy policies, but are loathe to refactor code to do so. Thus, Riverbed assumes that server-side code is weakly adversarial: poorly-designed applications may unintentionally try to leak data via explicit flows, but developers will not intentionally write code that attempts to surreptitiously leak data, e.g., via implicit flows, or via targeted attacks on the taint-tracking managed runtime.

A datacenter operator has physical access to servers, which enables direct manipulation of server RAM. So, our current Riverbed prototype assumes that datacenter operators are trusted. To ease this assumption, Riverbed could leverage a hardware-enforced isolation mechanism like Intel SGX [31]. SGX allows a secure computation to execute on a private core, with the hardware automatically encrypting and hashing the content of memory writes so that other cores are unable to inspect or modify the RAM that belongs to the secure computation. Ryoan [32] creates a distributed sandbox by representing an application as a directed graph of SGX computations that run on different servers. Unfortunately, SGX places limits on the memory size of secure applications. SGX also requires the applications to run in ring 3, forcing the code to rely on an untrusted OS in ring 0 to perform IO; the result is a large number of context switches for applications that perform many IOs [33]. Riverbed strives to be compatible with complex applications that often perform IO. So, Riverbed eschews mechanisms like SGX, and must be content to not protect against actively-malicious datacenter operators. To implement remote attestation, Riverbed does rely on trusted server-side TPM hardware, but TPMs do not protect against physical attacks on the main server hardware.

Riverbed assumes that the entire client-side is trusted, with the exception of the web content in a particular page. Buggy or malicious content may try to disclose too much information to a server. However, Riverbed ensures that whatever data is sent will be properly tagged. Since Riverbed uses TLS to authenticate network endpoints, the HTTPS certificate infrastructure must be trusted.

On the server, Riverbed’s TCB consists of the taint-tracking managed runtime, a reverse proxy that forwards requests to the appropriate universes, the TPM hardware [34] that provides the root of trust for attestation, and the daemon which servers use to attest to clients. We make standard cryptographic assumptions about the strength of the ciphers and hash functions used by the attestation protocol. Between the trusted hardware and the managed runtime are a boot loader, an OS, and possibly other systems software. Each end-user can select a different set of systems software to trust. These selections reside within the user’s policies, so that Riverbed’s client-side proxy can refuse to disclose data to untrusted server-side systems code.

4 Veil

4.1 Motivation

Web browsers are the client-side execution platform for a variety of online services. Many of these services handle sensitive personal data like emails and financial transactions. Since a user’s machine may be shared with other people, she may wish to establish a *private session* with a web site, such that the session leaves no persistent client-side state that can later be examined by a third party. Even if a site does not handle personally identifiable information, users may not want to leave evidence that a site was even visited. Thus, all popular browsers implement a private browsing mode which tries to remove artifacts like entries in the browser’s “recently visited” URL list.

Unfortunately, implementations of private browsing mode still allow sensitive information to leak into persistent storage [35, 36, 37, 38]. Browsers use the file system or a SQLite database to temporarily store information associated with private sessions; this data is often incompletely deleted and zeroed-out when a private session terminates, allowing attackers to extract images and URLs from the session. During a private session, web page state can also be reflected from RAM into swap files and hibernation files; this state is in cleartext, and therefore easily analyzed by curious individuals who control a user’s machine after her

private browsing session has ended. Simple greps for keywords are often sufficient to reveal sensitive data [35, 36].

Web browsers are complicated platforms that are continually adding new features (and thus new ways for private information to leak). As a result, it is difficult to implement even seemingly straightforward approaches for strengthening a browser’s implementation of incognito modes. For example, to prevent secrets in RAM from paging to disk, the browser could use OS interfaces like `mlock()` to pin memory pages. However, pinning may interfere in subtle ways with other memory-related functionality like garbage collecting or tab discarding [39]. Furthermore, the browser would have to use `mlock()` indiscriminately, on *all* of the RAM state belonging to a private session, because the browser would have no way to determine which state in the session is actually sensitive, and which state can be safely exposed to the swap device.

In this thesis, we introduce Veil, a system that allows web developers to implement private browsing semantics for their own pages. For example, the developers of a whistleblowing site can use Veil to reduce the likelihood that employers can find evidence of visits to the site on workplace machines. Veil’s privacy-preserving mechanisms are enforced *without assistance from the browser*—even if users visit pages using a browser’s built-in privacy mode, Veil provides stronger assurances that can only emerge from an intentional composition of HTML, CSS, and JavaScript. Veil leverages five techniques to protect privacy: URL blinding, content mutation, heap walking, DOM hiding, and state encryption.

- Developers pass their HTML and CSS files through Veil’s compiler. The compiler locates cleartext URLs in the content, and transforms those raw URLs into *blinded references* that are derived from a user’s secret key and are cryptographically unlinkable to the original URLs. The compiler also injects a runtime library into each page; this library interposes on dynamic content fetches (e.g., via `XMLHttpRequests`), and forces those requests to also use blinded references.
- The compiler uploads the objects in a web page to Veil’s *blinding servers*. A user’s

browser downloads content from those blinding servers, and the servers collaborate with a page’s JavaScript code to implement the blinded URL protocol. To protect the client-side memory artifacts belonging to a page, the blinding servers also *dynamically mutate* the HTML, CSS, and JavaScript in a page. Whenever a user fetches a page, the blinding servers create syntactically different (but semantically equivalent) versions of the page’s content. This ensures that two different users of a page will each receive unique client-side representations of that page.

- Ideally, sensitive memory artifacts would never swap out in the first place. Veil allows developers to mark JavaScript state and renderer state as sensitive. Veil’s compiler injects *heap walking code* to keep that state from swapping out. The code uses JavaScript reflection and forced DOM layouts to periodically touch the memory pages that contain secret data. This coerces the OS’s least-recently-used algorithm to keep the sensitive RAM pages in memory.
- Veil sites which desire the highest level of privacy can opt to use Veil’s *DOM hiding* mode. In this mode, the client browser essentially acts as a dumb graphical terminal. Pages are rendered on a content provider’s machine, with the browser sending user inputs to the machine via the blinding servers; the content provider’s machine responds with new bitmaps that represent the updated view of the page. In DOM hiding mode, the page’s unique HTML, CSS, and JavaScript content is never transmitted to the client browser.
- Veil also lets a page store private, persistent state by *encrypting* that state and by naming it with a blinded reference that only the user can generate.

By using blinded references for all content names (including those of top-level web pages), Veil avoids information leakage via client-side, name-centric interfaces like the DNS cache [40], the browser cache, and the browser’s address bar. Encryption allows a Veil page to safely leverage the browser cache to reduce page load times, or store user data across different sessions of

Browsing mode	Persistent, per-site client-side storage	Information leaks through client-side, name-based interfaces	Per-site browser RAM artifacts	GUI interactions
Regular browsing	Yes (cleartext by default)	Yes	Yes	Locally processed
Regular incognito mode	No	Yes	Yes	Locally processed
Veil with encrypted client-side storage, mutated DOM content, heap walking	Yes (always encrypted)	No (blinding servers)	Yes (but mutated and heap-walked)	Locally processed
Veil with DOM hiding	No	No (blinding servers)	No	Remotely processed

Table 1: A comparison between Veil’s two browsing modes, regular incognito browsing, and regular browsing that does not use incognito mode.

the private web page. A page that desires the highest level of security will eschew even the encrypted cache, and use DOM hiding; in concert with URL blinding, the hiding of DOM content means that the page will generate *no greppable state in RAM or persistent storage* that could later be used to identify the page. Table 1 summarizes the different properties of Veil’s two modes for private browsing.

In summary, **Veil is the first web framework that allows developers to implement privacy-preserving browsing semantics for their own pages.** These semantics are stronger than those provided by native in-browser incognito modes; however, Veil pages load on commodity browsers, and do not require users to reconfigure their systems or run their browsers within a special virtual machine [41]. Veil can translate legacy pages to more secure versions automatically, or with minimal developer assistance, easing the barrier to deploying privacy-preserving sites. Experiments show that Veil’s overheads are moderate: 1.25x–3.25x for Veil with encrypted client-side storage, mutated DOM content, and heap walking; and 1.2x–2.1x for Veil in DOM hiding mode.

4.2 Veil Overview

As shown in Figure 1, the Veil framework consists of three components. The *compiler* transforms a normal web page into a new version that implements static and dynamic privacy protections. Web developers upload the compiler’s output to *blinding servers*. These servers act as intermediaries between content publishers and content users, mutating and encrypting content. To load the Veil page, a user first loads a small *bootstrap page*. The bootstrap asks for a per-user key and the URL of the Veil page to load; the bootstrap then downloads

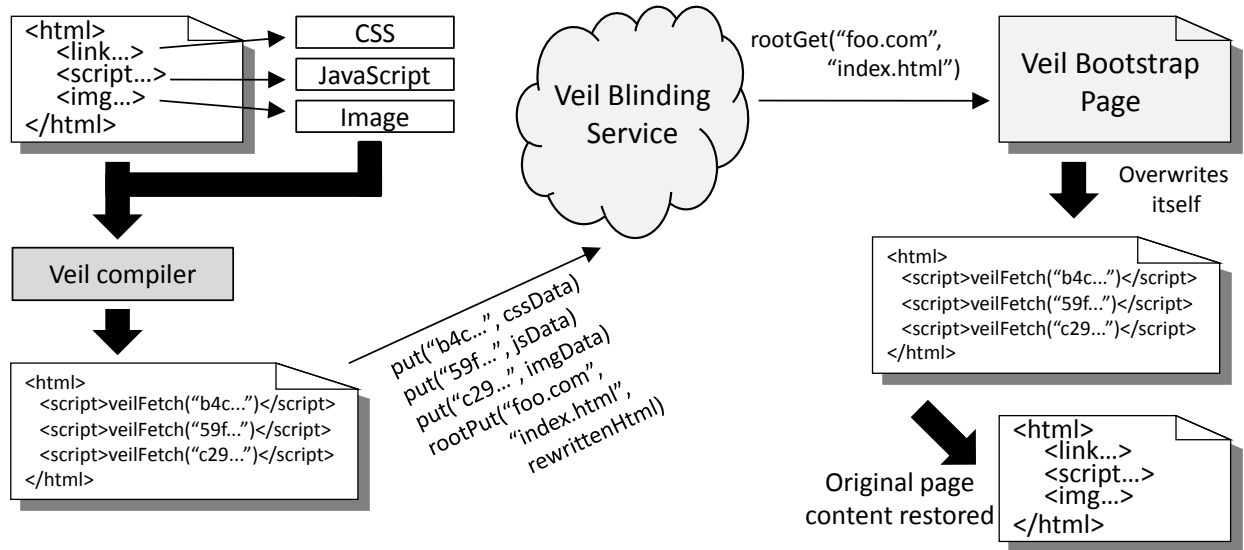


Figure 3: The Veil architecture (cryptographic operations omitted for clarity).

the appropriate objects from the blinding servers and dynamically overwrites itself with the privacy-preserving content in the target page.

4.3 Deployment Model

Veil uses an opt-in model, and is intended for web sites that want to actively protect client-side user privacy. For example, a whistleblowing site like SecureDrop [42] is incentivized to hide client-side evidence that the SecureDrop website has been visited; strong private browsing protections give people confidence that visiting SecureDrop on a work machine will not lead to incriminating aftereffects. As another example of a site that is well-suited for Veil, consider a web page that allows teenagers to find mental health services. Teenagers who browse the web on their parents' machines will desire strong guarantees that the machines store no persistent records of private browsing activity.

Participating Veil sites must explicitly recompile their content using the Veil compiler. This requirement is not unduly burdensome, since all non-trivial frameworks for web development impose a developer-side workflow discipline. For example, Aurelia [43], CoffeeScript [44], and Meteor [45] typically require a compilation pass before content can go live.

Participating Veil sites must also explicitly serve their content from Veil blinding servers.

Like Tor servers [46], Veil’s blinding servers can be run by volunteers, although content providers can also contribute physical machines or VMs to the blinding pool.

Today, many sites are indifferent towards the privacy implications of web browsing; other sites are interested in protecting privacy, but lack the technical skill to do so; and others are actively invested in using technology to hide sensitive user data. Veil targets the latter two groups of site operators. Those groups are currently in the minority, but they are growing. An increasing number of web services define their value in terms of privacy protections [47, 48, 49, 50], and recent events have increased popular awareness of privacy issues [51]. Thus, we believe that frameworks like Veil will become more prevalent as users demand more privacy, and site operators demand more tools to build privacy-respecting systems.

4.4 Threat Model

Veil assumes that a web service is actively interested in preserving its users’ client-side privacy. Thus, Veil trusts web developers and the blinding servers. Veil’s goal is to defend the user against local attackers who take control of a user’s machine *after* a private session terminates. If an attacker has access to the machine *during* a private session, the attacker can directly extract sensitive data, e.g., via keystroke logging or by causing the browser to core dump; such exploits are out-of-scope for this paper.

Veil models the post-session attacker as a skilled system administrator who knows the location and purpose of the swap file, the browser cache, and files like `/var/log/*` that record network activity like DNS resolution requests. Such an attacker can use tools like `grep` or `find` to look for hostnames, file types, or page content that was accessed during a Veil session. The attacker may also possess off-the-shelf forensics tools like Mandiant Redline [52] that look for traces of private browsing activity. However, the attacker lacks the skills to perform a customized, low-level forensics investigation that, e.g., tries to manually extract C++ data structures from browser memory pages that Veil could not prevent from swapping out.

Given this attacker model, Veil’s security goals are weaker than strict forensic deniability [41]. However, Veil’s weaker type of forensic resistance is both practically useful and,

in many cases, the strongest guarantee that can be provided without forcing users to run browsers within special OSes or virtual machines. Veil’s goal is to load pages within *unmodified* browsers that run atop *unmodified* operating systems. Thus, Veil is forced to implement privacy-preserving features using browser and OS interfaces that are unaware of Veil’s privacy goals. These constraints make it impossible for Veil to provide strict forensic deniability. However, most post-session attackers (e.g., friends, or system administrators at work, Internet cafes, or libraries) will lack the technical expertise to launch FBI-style forensic investigations.

Using blinded URLs, Veil tries to prevent data leaks through system interfaces that use network names. Examples of name-based interfaces are the browser’s “visited pages” history, the browser cache, cookies, and the DNS cache (which leaks the hostnames of the web servers that a browser contacts [35]). It is acceptable for the attacker to learn that a user has contacted Veil’s blinding servers—those servers form a large pool whose hostnames are generic (e.g., `veil.io`) and do not reveal any information about particular Veil sites.

Veil assumes that web developers only include trusted content that has gone through the Veil compiler. A page may embed third party content like a JavaScript library, but the Veil compiler analyzes both first party and third party content during compilation.

Heap walking allows Veil to prevent sensitive RAM artifacts from swapping to disk. Veil does not try to stop information leaks from GPU RAM [53], but GPU RAM is never swapped to persistent storage. Poorly-written or malicious browser extensions that leak sensitive page data [54] are also outside the scope of this paper.

5 Proposed Timeline

Modify text for Riverbed system (4 weeks, finish by end of May)

Schedule and do PhD defense (finish by end of May)

Take papers and typeset them into PhD thesis format (2 days)

Write a new introduction that better encompasses all systems (3 days)

Fix bibliography (1 day)

Write acknowledgments (1 day)

Final fixes on formatting and figures (2 days)

Submit thesis by mid-June

References

- [1] Arvind Narayanan and Vitaly Shmatikov. Myths and fallacies of personally identifiable information. *Communications of the ACM*, 53(6):24–26, 2010.
- [2] Arvind Narayanan and Vitaly Shmatikov. Robust de-anonymization of large sparse datasets. In *Proceedings of the 29th IEEE Symposium on Security and Privacy*, pages 111–125, Oakland, CA, May 2008.
- [3] Adam Tanner. Different customers, Different prices, Thanks to big data, March 21 2014. <http://www.forbes.com/sites/adamtanner/2014/03/26/different-customers-different-prices-thanks-to-big-data/>.
- [4] Felix Salmon. Why the Internet is Perfect for Price Discrimination, September 3 2013. <http://blogs.reuters.com/felix-salmon/2013/09/03/why-the-internet-is-perfect-for-price-discrimination/>.
- [5] Aniko Hannak, Gary Soeller, David Lazer, Alan Mislove, and Christo Wilson. Measuring price discrimination and steering on e-commerce web sites. In *Proceedings of the Conference on Internet Measurement Conference (IMC)*, pages 305–318, 2014.
- [6] Jeremy Singer-Vine Jennifer Valentino-Devries and Ashkan Soltani. Websites Vary Prices, Deals Based on Users’ Information, December 24 2012. Wall Street Journal.
- [7] Rick Seaney. Do Cookies Really Raise Airfares?, April 30 2013. <http://www.usatoday.com/story/travel/columnist/seaney/2013/04/30/airfare-expert-do-cookies-really-raise-airfares/2121981/>.

- [8] Ramprasad Ravichandran, Michael Benisch, Patrick Gage Kelley, and Norman M Sadeh. Capturing social networking privacy preferences. In *Proceedings of the 9th Privacy Enhancing Technologies Symposium*, pages 1–18, Seattle, WA, August 2009.
- [9] Jason Kincaid. Another Security Hole Found on Yelp, Facebook Data Once Again Put at Risk, May 11 2010. <http://techcrunch.com/2010/05/11/another-security-hole-found-on-yelp-facebook-data-once-again-put-at-risk/>.
- [10] Fahmida Y. Rashid. Twitter Breached, Attackers Stole 250,000 User Data, February 2 2013. <http://securitywatch.pcmag.com/none/307708-twitter-breached-attackers-stole-250-000-user-data>.
- [11] Femi Olumofin and Ian Goldberg. Privacy-preserving queries over relational databases. In *Proceedings of the 10th Privacy Enhancing Technologies Symposium*, pages 75–92, Berlin, Germany, 2010.
- [12] Benny Chor, Niv Gilboa, and Moni Naor. *Private information retrieval by keywords*. Citeseer, 1997.
- [13] Joel Reardon, Jeffrey Pound, and Ian Goldberg. Relational-complete private information retrieval. *University of Waterloo, Tech. Rep. CACR*, 34:2007, 2007.
- [14] David J Wu, Joe Zimmerman, J  r  my Planul, and John C Mitchell. Privacy-preserving shortest path computation. In *Proceedings of the 2016 Annual Network and Distributed System Security Symposium*, San Diego, CA, February 2016.
- [15] Chang Lan, Justine Sherry, Raluca Ada Popa, Sylvia Ratnasamy, and Zhi Liu. Embark: Securely outsourcing middleboxes to the cloud. In *Proceedings of the 13th Symposium on Networked Systems Design and Implementation (NSDI)*, pages 255–273, Santa Clara, CA, March 2016.

- [16] Assaf Ben-David, Noam Nisan, and Benny Pinkas. FairplayMP: A system for secure multi-party computation. In *Proceedings of the 15th ACM Conference on Computer and Communications Security (CCS)*, pages 257–266, Alexandria, VA, October 2008.
- [17] Elette Boyle, Niv Gilboa, and Yuval Ishai. Function secret sharing. In *Proceedings of the 34th Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, pages 337–367. Sofia, Bulgaria, April 2015.
- [18] Niv Gilboa and Yuval Ishai. Distributed point functions and their applications. In *Proceedings of the 33rd Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, pages 640–658. Copenhagen, Denmark, May 2014.
- [19] Jeffrey Rott. Intel advanced encryption standard instructions (AES-NI). Technical report, Technical report, Intel, 2010.
- [20] OpenStreetMap. OpenStreetMap. <https://www.openstreetmap.org/>.
- [21] Renu Tewari, Thirumale Niranjan, and Srikanth Ramamurthy. WCDP: A protocol for web cache consistency. In *Proceedings of the 7th Web Caching Workshop*. Citeseer, 2002.
- [22] Chi-Hung Chi, Choon-Keng Chua, and Weihong Song. A novel ownership scheme to maintain web content consistency. In *International Conference on Grid and Pervasive Computing*, pages 352–363. Springer, 2008.
- [23] Pascal Paillier. Public-key cryptosystems based on composite degree residuosity classes. In *Proceedings of the 18th Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, pages 223–238, Prague, Czech Republic, May 1999.
- [24] A. Booth. Charities hit with fines for sharing donors’ data without consent, December 7, 2016. Sophos Naked Security Blog. <https://nakedsecurity.sophos.com/2016/12/07/charities-hit-with-fines-for-sharing-donors-data-without-consent/>.

- [25] P. Sayer. German consumer groups sue WhatsApp over privacy policy changes, January 30, 2017. PCWorld. <http://www.pcworld.com/article/3163027/private-cloud/german-consumer-groups-sue-whatsapp-over-privacy-policy-changes.html>.
- [26] K. Zetter. Hackers Finally Post Stolen Ashley Madison Data, August 18, 2015. Wired. <https://www.wired.com/2015/08/happened-hackers-posted-stolen-ashley-madison-data/>.
- [27] EU Parliament. GDPR Portal, 2017. <http://www.eugdpr.org/eugdpr.org.html>.
- [28] D. Hedin and A. Sabelfeld. A Perspective on Information-Flow Control. In *Proceedings of the 2011 Marktoberdorf Summer School*, August 2001.
- [29] P. Li, Y. Mao, and S. Zdancewic. Information Integrity Policies. In *Proceedings of the Workshop on Formal Aspects in Security and Trust*, September 2003.
- [30] G. Coker, J. Guttman, P. Loscocco, A. Herzog, J. Millen, B. O’Hanlon, J. Ramsdell, A. Segall, and B. Sniffen. Principles of remote attestation. *International Journal of Information Security*, 10(2):63–81, June 2011.
- [31] V. Costan and S. Devadas. Intel SGX Explained, February 20, 2017. Cryptology ePrint Archive: Version 20170221:054353. <https://eprint.iacr.org/2016/086.pdf>.
- [32] T. Hunt, Z. Zhu, Y. Xu, S. Peter, and E. Witchel. Ryoan: A Distributed Sandbox for Untrusted Computation on Secret Data. In *Proceedings of OSDI*, pages 533–549, 2016.
- [33] A. Baumann, M. Peinado, and G. Hunt. Shielding Applications from an Untrusted Cloud with Haven. In *Proceedings of OSDI*, pages 267–283, 2014.
- [34] Trusted Computing Group. Trusted Platform Module (TPM) Summary. <https://trustedcomputinggroup.org/trusted-platform-module-tpm-summary/>.

- [35] G. Aggarwal, E. Burzstein, C. Jackson, and D. Boneh. An Analysis of Private Browsing Modes in Modern Browsers. In *Proceedings of USENIX Security*, Washington, DC, August 2010.
- [36] D. Isacsson. Microsoft Edge’s Incognito Mode Isn’t So Incognito, February 1, 2016. Digital Trends. <https://www.digitaltrends.com/web/microsoft-edge-security-flaws-in-incognito/>.
- [37] Magnet Forensics. How Does Chrome’s ”Incognito” Mode Affect Digital Forensics? <http://www.magnetforensics.com/how-does-chromes-incognito-mode-affect-digital-forensics/>, August 6, 2013.
- [38] D. Ohana and N. Shashidhar. Do Private and Portable Web Browsers Leave Incriminating Evidence? In *Proceedings of the International Workshop on Cyber Crime*, San Francisco, CA, May 2013.
- [39] A. Osmani. Tab Discarding in Chrome: A Memory-Saving Experiment, September 2015. Google Developer Blog. <https://developers.google.com/web/updates/2015/09/tab-discarding>.
- [40] E. Felten and M. Schneider. Timing Attacks on Web Privacy. In *Proceedings of CCS*, Athens, Greece, November 2000.
- [41] A. Dunn, M. Lee, S. Jana, S. Kim, M. Silberstein, Y. Xu, V. Shmatikov, and E. Witchel. Eternal Sunshine of the Spotless Machine: Protecting Privacy with Ephemeral Channels. In *Proceedings of OSDI*, Vancouver, BC, Canada, November 2010.
- [42] New Yorker. The New Yorker SecureDrop, 2017. <https://projects.newyorker.com/securedrop/>.
- [43] Blue Spire Inc. Aurelia, 2017. <http://aurelia.io/>.

- [44] CoffeeScript. CoffeeScript: A Little Language that Compiles to JavaScript, October 26, 2017. <http://coffeescript.org/>.
- [45] Meteor Development Group. Meteor: The Fastest Way to Build JavaScript Apps, 2017. <https://www.meteor.com/>.
- [46] R. Dingledine, N. Mathewson, and P. Syverson. Tor: The Second-Generation Onion Router. In *Proceedings of USENIX Security*, San Diego, CA, August 2004.
- [47] DuckDuckGo. Take back your privacy! Switch to the search engine that doesn't track you., 2017. <https://duckduckgo.com/about>.
- [48] Enigma. Arrival Data for Non-Stop Domestic Flights by Major Air Carriers for 2012. <https://app.enigma.io/table/us.gov.dot.rita.trans-stats.on-time-performance.2012>.
- [49] Priv.io. Priv.io homepage, 2015. <https://priv.io/>.
- [50] Priv.ly. Change the Way Your Browser Works: Share Priv(ate).ly, 2017. <https://priv.ly/>.
- [51] E. Orion. Tor popularity leaps after snooping revelations, August 30, 2013. The Inquirer. <http://www.theinquirer.net/inquirer/news/2291758/tor-popularity-leaps-after-snooping-revelations>.
- [52] Mandiant. Mandiant Redline Users Guide, 2012. https://dl.mandiant.com/EE/library/Redline1.7_UserGuide.pdf.
- [53] S. Lee, Y. Kim, J. Kim, and J. Kim. Stealing Webpages Rendered on Your Browser by Exploiting GPU Vulnerabilities. In *Proceedings of IEEE Symposium on Security and Privacy*, San Jose, CA, May 2014.

- [54] B. Lerner, L. Elberty, N. Poole, and S. Krishnamurthi. Verifying Web Browser Extensions' Compliance with Private-Browsing Mode. In *Proceedings of ESORICS*, Egham, United Kingdom, September 2013.